

# Homework 1

```
In [3]: ###module import
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from lmfit import Model
from lmfit import Parameters
from scipy.optimize import minimize
from scipy.optimize import fmin_tnc
from sklearn.model_selection import KFold
import seaborn as sns
from sklearn.linear_model import LogisticRegression, RidgeClassifier, Lasso
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV, StratifiedKFold
from sklearn.metrics import roc_auc_score, fl_score, roc_curve
```

## Q1

```

In [4]: ### load the data
train=pd.DataFrame([[0,1],[2,4],[3,9],[5,16]])
test=pd.DataFrame([[1,3],[4,12]])

x=train.iloc[:,0].values
y=train.iloc[:,1].values

###define the model
#ref for the following: https://stackoverflow.com/questions/48469889/how-to-fit-a-polynomi
color_for_degree=["purple","blue","cyan","green","orange"]
coef_name=['a0','a1','a2','a3','a4']

def func(x,a0,a1,a2,a3,a4):
    return a0 + a1*x + a2*x ** 2 + a3*x**3 + a4*x**4

#fit the model -> plot

###plot the training and the test points
plt.plot(x, y, 'bo',color="blue")
plt.plot(test.iloc[:,0].values, test.iloc[:,1].values, 'bo',color="red")

#I am able to wrap the following into a for loop by this: https://lmfit.github.io/lmfit-py

ds=[4,3,2,1,0]
coef_name=['a0','a1','a2','a3','a4']

pmodel = Model(func)
params = Parameters()
results=[]

for d1 in ds:
    for d2 in ds:
        if d2>d1:
            params.add(coef_name[d2],0)
            params[coef_name[d2]].vary = False
        elif d2==d1:
            params.add(coef_name[d2],1)
            params[coef_name[d2]].vary = False
        else:
            params.add(coef_name[d2],1)

    result = pmodel.fit(y, params, x=x)
    results.append(result)

    xnew = np.linspace(x[0], x[-1], 1000)
    ynew = result.eval(x=xnew)

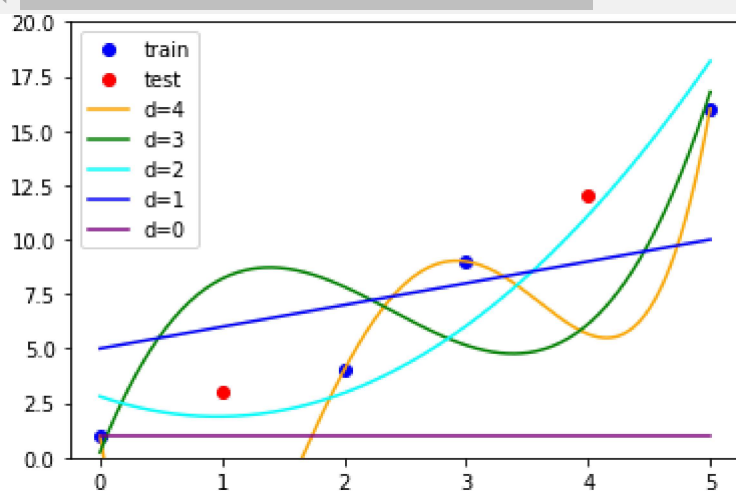
    plt.ylim(bottom=0, top=20)
    plt.plot(xnew, ynew, 'r-',color=color_for_degree[d1])

```

```
### show everything
```

```
plt.gca().legend(('train','test','d=4','d=3','d=2','d=1','d=0')) #ref: https://stackoverflow.com/questions/10114244/matplotlib-legend-order
```

```
plt.show()
```



In [5]: #Ref: <https://datascienceomar.wordpress.com/2016/07/03/bias-and-variance-with-scikit-learn>

```
x=np.array([0,1,2,3,4,5])
y=np.array([1,3,4,9,12,16])

bias=[]
variance=[]
total_error=[]
train_error=[]
test_error=[]

for d in reversed(ds):
    ##bias = E(f_hat(x)) - E(f(x)) = mean(f_hat(x) - mean(y))
    y_hat=results[d].eval(x=x)

    bias.append( np.mean(y_hat-np.mean(y)) )

    ##variance=E(f_hat(x)**2)-E(f_hat(x))**2. I don't have the intuitive understanding of
    variance.append( np.mean(y_hat**2)-np.mean(y_hat)**2 )

    ##total error rate = traning error + test error? How do we define test error?
    ##(?)error=RSS


    total_error.append( sum( (y-y_hat)**2 ) )

    train_error.append( sum( (y-y_hat)**2 )[np.array([0,2,3,5])] ) )

    test_error.append( sum( (y-y_hat)**2 )[np.array([1,4])] ) )

for metric in [bias,variance,total_error,train_error,test_error]:
    plt.scatter(ds,metric)
    plt.plot(ds,metric)

plt.xlabel('degree')
plt.gca().legend(('bias','variance','total_error','train_error','test_error')) #ref: https
plt.show()
```



At  $d=2$  it seems to have the best fit because the 5 metrics (bias, variance, total error, training error, test error) are generally smaller than in other degrees. At lower degrees ( $d=0,1$ ) the models seem to underfit, while on higher degrees ( $d=3,4$ ) the models seem to overfit (high variance)

```

In [6]: smarket=pd.read_csv("Smarket.csv")
y=smarket['Today'].values
X=smarket[['Lag1','Lag2']].values

def model_func(X,params):
    # 2d Line Z = aX + b
    return X.dot(params[:2]) + params[2]

##have to modify the following code

for p in [1, 2]:

    L2_norm=[]

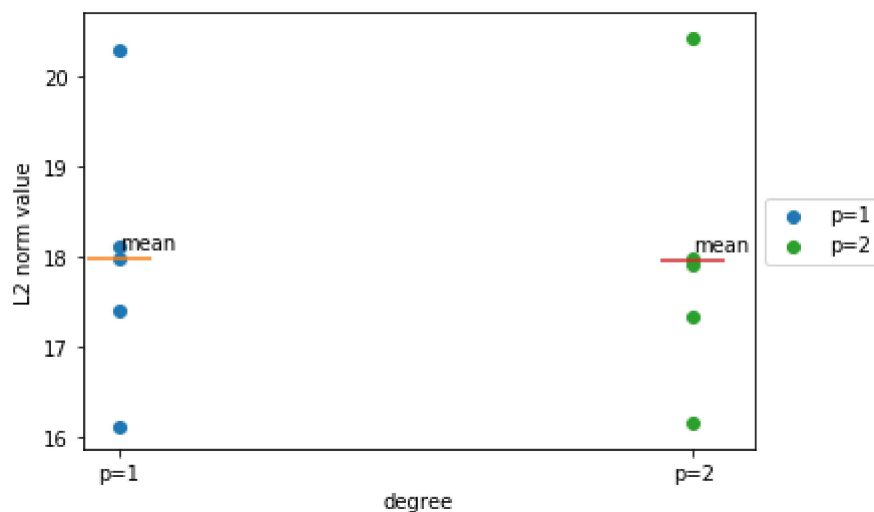
    def cost_function(params, X, y, p):
        error_vector = y - model_func(X, params)
        return np.linalg.norm(error_vector, ord=p) #this guy returns the Lp norm

    kf=KFold(n_splits=5, random_state=101, shuffle=True)
    for train_index, test_index in kf.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]
        output = minimize(cost_function, [1,1,1], args=(X_train, y_train, p)) #[1,1,1] is
        L2_norm.append( np.sqrt( sum ( (y_test-model_func(X_test, output.x))**2 ) ) ) #L2

    label='p='+str(p)
    plt.scatter([label]*5,L2_norm,label=label)
    mean=np.mean(L2_norm)
    plt.scatter([label],mean,s=1000,marker="_") #s means size
    plt.annotate('mean',xy=(p-1,mean+0.1))

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.xlabel('degree')
plt.ylabel('L2 norm value')
plt.show()

```



Using L1 or L2 as the minimization method seems to give similar result based on the assesment on the L2 norm value. This is likely because no matter how we change p in Lp norm, the problem stays the same as a linear regression question. Therefore, we don't get significant performace boost by changing the norms

## Q2

In [ ]:

$$\text{cross entropy error} = - \sum \{ y_i \log [\sigma(x + \beta_0)] + (1 - y_i) \log [1 - \sigma(x + \beta_0)] \}$$

(1)

$$\sum_{n=1}^2 (\sigma(x + \beta_0) - y_n) x_n = 0$$

$$[\sigma(-3 + \beta_0) - 1] \times (-3) +$$

$$[\sigma(-1 + \beta_0) - 0] \times (-1) = 0$$

$$\sigma = \text{sigmoid}() = \frac{1}{1 + e^{-x}}$$

(2)

$$\text{model: } p(x) = \frac{e^{x + \beta_0}}{1 + e^{x + \beta_0}} = \frac{e^{x + 3.78}}{1 + e^{x + 3.78}}$$

$$p(-4) = \frac{e^{-0.22}}{1 + e^{-0.22}} = \underline{0.475} \text{ Answer}$$

$$p(5) = \frac{e^{8.78}}{1 + e^{8.78}} = \underline{0.9998} \text{ Answer}$$

$$\boxed{e = \beta_1 \quad e^{-\beta_0} = x}$$

$$\frac{-3}{1 + \beta_1^3 x} + 3 - \frac{1}{1 + \beta_1 x} = 0$$

$$-3(1 + \beta_1 x) - (1 + \beta_1^3 x) = -3(1 + \beta_1 x)(1 + \beta_1^3 x)$$

$$-4 - 3\beta_1 x - \beta_1^3 x = -3 - 3(\beta_1 + \beta_1^3)x - 3\beta_1^4 x^2$$

$$3\beta_1^4 x^2 + (-3\beta_1 - \beta_1^3 + 3\beta_1 + 3\beta_1^3)x - 1 = 0$$

$$3\beta_1^4 x^2 + 2\beta_1^3 x - 1 = 0$$

$$3 \cdot e^4 x^2 + 2 \cdot e^3 x - 1 = 0$$

$$163.79445 x^2 + 40.171 x - 1 = 0$$

$$x = -0.26, 0.022778$$

$$\beta_0 = -\ln x$$

$$= -\ln 0.022778$$

$$= \underline{3.78196} \text{ Answer}$$

Q3




```
In [9]: hwl_input=pd.read_csv("hw1_input.csv")
hw1_input.dtypes #14 variables
hw1_input.describe()

new_hwl_input=pd.get_dummies(hwl_input)
new_hwl_input.shape
new_hwl_input.dtypes

sns.pairplot(new_hwl_input) #scatter plot + histogram
```

Out[9]: <seaborn.axisgrid.PairGrid at 0x1e65a3f87b8>



From the above we know:

- (1) Age, BP, Cholestoral, max hr, major vessels are discrete variables (integer).
- (2) oldpeak is a continuous variable
- (3) Others columns are categorical variables transformed into multiple dummy variables

```

In [39]: import warnings
warnings.filterwarnings('ignore') #I have made sure that all following warnings don't matter

X = new_hwl_input.iloc[:, 0:-2]
y = hwl_input.iloc[:, -1]
y[y=='Yes']=1
y[y=='No']=0
y=y.astype('int')

ridge_coef=[]
lasso_coef=[]
ridgeAUROC=[]
lassoAUROC=[]
best_ridge_F1=[]
best_lasso_F1=[]
best_ridge_F1_threshold=[]
best_lasso_F1_threshold=[]

for seed in range(1000): #change to 1000 after debugging
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=seed)

    #for Ridge

    ridge_classifier = RidgeClassifier()
    ridge_classifier.fit(X_train, y_train.values) #ref: https://www.geeksforgeeks.org/num
    ridge_coef.append(ridge_classifier.coef_[0])
    class_=ridge_classifier.predict(X_test)

    #====In this part I manually calculated the probability because I haven't found a way
    beta_X=np.dot(ridge_classifier.coef_,np.transpose(X_test))
    probas=np.e**beta_X/(1+np.e**beta_X)
    probas=np.transpose(probas[0].tolist())
    #====

    ridge_F1=[]
    for threshold in probas:
        y_pred=(probas>threshold).astype(int)
        ridge_F1.append(f1_score(y_test, y_pred))

    best_ridge_F1.append(max(ridge_F1))
    best_ridge_F1_threshold.append(probas[ridge_F1.index(max(ridge_F1))])
    ridgeAUROC.append(roc_auc_score(y_test, ridge_classifier.predict(X_test)))

    #get ROC parameter
    if seed==0:
        fpr1, tpr1, thresholds1=roc_curve(y_test,probas)
        auroc1=roc_auc_score(y_test, ridge_classifier.predict(X_test))

    #for lasso (not sure how it knows when to use linear regression or logistic regression)
    lasso = Lasso()

```

```
lasso.fit(X_train, y_train)
lasso_coef.append(lasso.coef_)

probas=lasso.predict(X_test)
probas=np.transpose(probas.tolist())

lasso_F1=[]
for threshold in probas:
    y_pred=(probas>threshold).astype(int)
    lasso_F1.append(f1_score(y_test, y_pred))

best_lasso_F1.append(max(lasso_F1))
best_lasso_F1_threshold.append(probas[lasso_F1.index(max(lasso_F1))])
lassoAUROC.append(roc_auc_score(y_test, lasso.predict(X_test)))

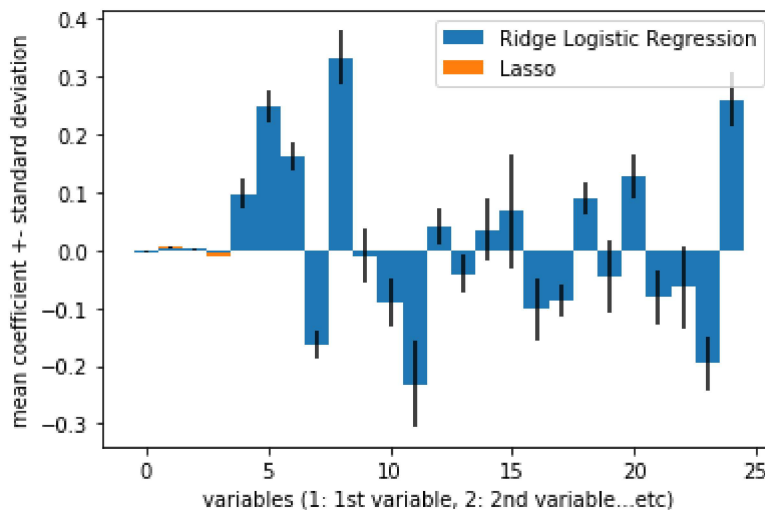
#get ROC parameter
if seed==0:
    fpr2, tpr2, thresholds2=roc_curve(y_test,probas)
    auROC2=roc_auc_score(y_test, lasso.predict(X_test))
```

## Coefficients

```
In [42]: avg_coeff_Ridge = np.mean(ridge_coef, axis = 0).tolist()
std_coeff_Ridge = np.std(ridge_coef, axis = 0).tolist()

avg_coeff_Lasso = np.mean(lasso_coef, axis = 0).tolist()
std_coeff_Lasso = np.std(lasso_coef, axis = 0).tolist()

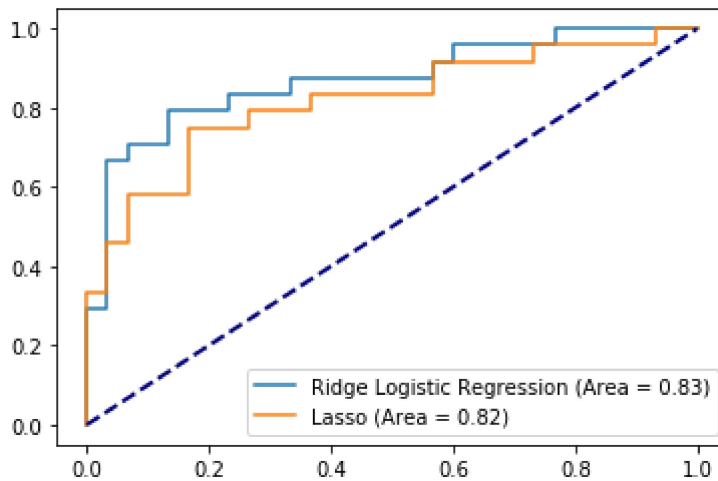
#ref for the following bargraph: https://matplotlib.org/3.1.1/gallery/lines_bars_and_markers
N = X_train.shape[1]
width = 1 # the width of the bars: can also be len(x) sequence
ind = np.arange(N) # the x locations for the groups
p1 = plt.bar(ind, avg_coeff_Ridge, width, yerr=std_coeff_Ridge, label="Ridge Logistic Regression")
p2 = plt.bar(ind, avg_coeff_Lasso, width,
             bottom=avg_coeff_Ridge, yerr=std_coeff_Lasso, label="Lasso")
plt.legend()
plt.xlabel('variables (1: 1st variable, 2: 2nd variable...etc)')
plt.ylabel('mean coefficient +- standard deviation')
plt.show()
```



From the original data, the multi-class columns have been transformed to dummy variables (0,1) in order to implement the regression. It seems for Ridge most of the variables matter, but some coefficients have high standard deviation, indicating that the values of those variables are not that robust. For Lasso only 3 variables matter (BP, Cholesterol, max hr), yet the corresponding coefficients are relatively small compared to coefficients in Ridge.

## ROC curve

```
In [43]: plt.plot(fpr1, tpr1, label="Ridge Logistic Regression (Area = %0.2f)" % auroc1)
plt.plot(fpr2, tpr2, label="Lasso (Area = %0.2f)" % auroc2)
plt.legend()
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.show()
```



Area = area under the ROC curve (perfect = 1)

## Best threshold for F1

```
In [44]: print('Best F1 probability cutoff for Ridge: ', best_ridge_F1_threshold[best_ridge_F1.index])
print('Best F1 probability cutoff for Lasso: ', best_lasso_F1_threshold[best_lasso_F1.index])
```

Best F1 probability cutoff for Ridge: 0.6167453553701893

Best F1 probability cutoff for Lasso: 0.4452220390013467

## ### Mean + Standard deviation for AUROC

```
In [45]: print('Mean AUROC for Ridge +- Std: ', np.mean(ridgeAUROC), ' + ', np.std(best_ridge_F1))
print('Mean AUROC for Lasso +- Std: ', np.mean(lassoAUROC), ' + ', np.std(best_lasso_F1))
```

Mean AUROC for Ridge +- Std: 0.83705 + 0.04398081291667096

Mean AUROC for Lasso +- Std: 0.7411020833333335 + 0.04206013704000325

## ### Mean + Standard deviation for best F1

```
In [46]: print('Mean F1 score for Ridge +- Std: ', np.mean(best_ridge_F1), ' + ', np.std(best_ridge_F1))  
         print('Mean F1 score for Lasso +- Std: ', np.mean(best_lasso_F1), ' + ', np.std(best_lasso_F1))
```

```
Mean F1 score for Ridge +- Std: 0.8529650861163477 + 0.04398081291667096  
Mean F1 score for Lasso +- Std: 0.7121581524949887 + 0.04206013704000325
```

```
In [ ]:
```