

Homework 2

This jupyter notebook is also accessble as an online colab notebook :

<https://colab.research.google.com/drive/10KcuDSp9jCx0eFkdQ-UvrS6VPQjmpYPQ#scrollTo=Yj8OAlyA5Tzv>
(<https://colab.research.google.com/drive/10KcuDSp9jCx0eFkdQ-UvrS6VPQjmpYPQ#scrollTo=Yj8OAlyA5Tzv>)

Q1

Load & normalize the data

In [2]:

```
import pandas as pd
import numpy as np
from urllib.request import urlopen
import re
from sklearn import preprocessing

data_train = urlopen("https://raw.githubusercontent.com/peterwu19881230/CSCE633_ML_data/master/train.csv")
data_test = urlopen("https://github.com/peterwu19881230/CSCE633_ML_data/raw/master/test.csv")
#Note for the test data: 1. I manually deleted a "NUL" on line 1181 1. I manually deleted a "NUL" on line 1181

def parse_data(data):
    rows=[]
    class_=[]
    for line in data: # files are iterable

        string=line.decode("utf-8").rstrip() #rstrip() is to trim EOL
        words=string.split(sep=" ")

        class_.append(words[0])

        words.pop(0) #remove the first element (response variable)

        row=[np.nan]*36 #I have checked that for each row there are at most 36 features
        for word in words:
            col_result=re.search('([0-9]{1,2}):.*',word) #.* means "0 or more of any character"
            col=int(col_result.group(1))

            value_result=re.search('[0-9]{1,2}:(.*)',word)
```

```
row[col-1]=float(value_result.group(1))
```

```
rows.append(row)
```

```
df=pd.DataFrame(rows)
```

```
#normalize the data (definition of normalization: https://scikit-learn.org/stable/  
min_max_scaler = preprocessing.MinMaxScaler(feature_range=(0, 1))  
normalized_df=pd.DataFrame(min_max_scaler.fit_transform(df.T), columns=df.index, )  
X=normalized_df
```

```
##compute response variable
```

```
y=np.where(np.array(class_).astype(int) == 6, 1, 0)
```

```
return(X,y)
```

```
X_train, y_train = parse_data(data_train)
```

```
X_test, y_test = parse_data(data_test)
```

```
#check no. of cols & rows
```

```
print(X_train.head())
```


```
##verify by printing histograms for all columns
```

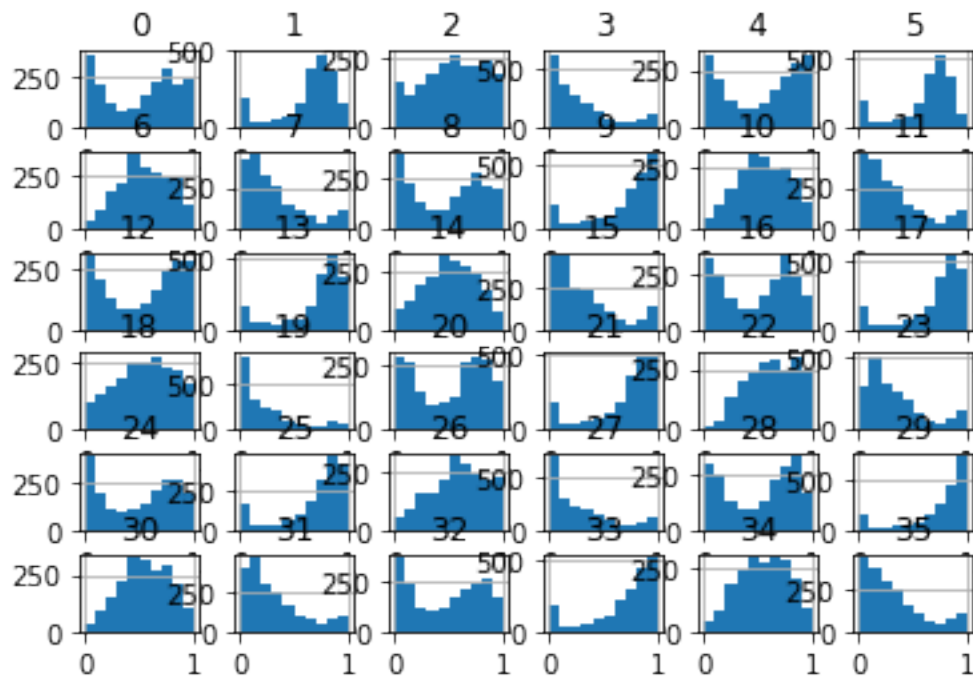

```
X_train.hist(); #";" is to surpress messages
```

```
X_test.hist();
```

```
0      1      2      3      ...      32      33      3
4      35
0  NaN  1.000000  0.694047  0.427273  ...  0.610156  0.651456  0.23157
9  0.158984
1  NaN  0.890850  0.196389  0.150320  ...  0.796417  0.839826  0.30722
7  0.141549
2  1.0  0.928588  0.771404  0.364708  ...  0.283042  0.448587  0.11340
0  0.002494
3  NaN  0.695489  0.229246  0.113419  ...  0.839346  0.831158  0.67827
0  0.437713
4  NaN  0.567578  0.280369  0.081376  ...  0.553549  0.451600  0.12377
6  0.000000
```

```
[5 rows x 36 columns]
```





Calculate how many nans are there -> Decide to ignore or impute -> Impute

In [3]:

```
print(X_train.isna().sum())

import numpy as np
from sklearn.impute import SimpleImputer

def impute_by_mean(data):
    imp_mean = SimpleImputer(missing_values=np.nan, strategy='mean') #It seems that t
    imp_mean.fit(data)
    return(pd.DataFrame(imp_mean.transform(data)))

X_train_imputed=impute_by_mean(X_train) #print(X_train_imputed.isna().sum()) #this
X_test_imputed=impute_by_mean(X_test)
```

```
0      59
1       0
2      55
3       0
4     231
5       0
6       0
7       0
8      54
9       0
```

10	0
11	0
12	259
13	230
14	0
15	172
16	55
17	0
18	0
19	0
20	0
21	0
22	0
23	0
24	63
25	219
26	0
27	156
28	0
29	0
30	0
31	0
32	59
33	0
34	0
35	0

dtype: int64

implement SVM

In [4]:

```
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import cross_val_score

#=====
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning) #avoid the "future wa
#=====

ds = [1,2,3,4]
Cs = [0.001, 0.01, 0.1, 1, 10] #ref: https://medium.com/@aneesha/svm-parameter-tuning
param_grid={'degree':ds, 'C':Cs}

grid_search = GridSearchCV(SVC(kernel='poly'), param_grid, cv=10,scoring='accuracy')
grid_search.fit(X_train_imputed, y_train)
```

Out[4]:

```
GridSearchCV(cv=10, error_score='raise-deprecating',
             estimator=SVC(C=1.0, cache_size=200, class_weight=None, c
             coef0=0.0,
                                decision_function_shape='ovr', degree=3,
                                gamma='auto_deprecated', kernel='poly', max
             _iter=-1,
                                probability=False, random_state=None, shrin
             king=True,
                                tol=0.001, verbose=False),
             iid='warn', n_jobs=None,
             param_grid={'C': [0.001, 0.01, 0.1, 1, 10],
                         'degree': [1, 2, 3, 4]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=F
             else,
             scoring='accuracy', verbose=0)
```

In [5]:

```
#I here use accuracy as the error term

accuracy_list=[]
for i in range(10):
    accuracy_list.append(pd.DataFrame(grid_search.cv_results_['split'+str(i)+'_test_sc



accuracy_df=pd.concat(accuracy_list,axis=0)



## assign column names
params=grid_search.cv_results_['params']
params_string=['C:'+str(param['C'])+'_'+str(param['degree'])] for param in
accuracy_df.columns=params_string
```

```
mean_array=np.array([np.mean(accuracy_df[column]) for column in accuracy_df.columns])
std_array=np.array([np.std(accuracy_df[column]) for column in accuracy_df.columns])
```

```
import matplotlib.pyplot as plt
```

```
for d in ds:
    degree_index=list(np.array([0,4,8,12,16])+d-1)
    plt.plot(Cs,mean_array[degree_index])
    plt.errorbar(Cs, mean_array[degree_index], yerr=std_array[degree_index],lolims=True)
    plt.scatter(Cs,mean_array[degree_index])
    plt.title('Degree = '+str(d))
    plt.xscale('log')
    plt.xlabel('C')
    plt.ylabel('Mean Accuracy +- std')
    plt.show()
```





Best parameters that gives best accuracy:

In [6]:

```
grid_search.best_params_
```

Out[6]:

```
{'C': 10, 'degree': 3}
```


In [7]:

```
from sklearn.metrics import accuracy_score

ds = [1,2,3,4]
C=10

accuracy_list=[]
no_support_vectors=[]
for d in ds:
    model=SVC(C=C,kernel='poly',degree=d)
    model.fit(X_train_imputed,y_train)
    no_support_vectors.append(model.n_support_) #no of support vectors for each feature
    accuracy_list.append(accuracy_score(y_test,model.predict(X_test_imputed)))

x=ds
plt.scatter(x,accuracy_list)
plt.plot(x,accuracy_list)
plt.xticks(np.arange(min(x), max(x)+1, 1.0))
plt.xlabel('degree')
plt.ylabel('Accuracy')
plt.show()
```



In [8]:


```
no_support_vectors #no. of support vectors for each class (for d=1,2,3,4, respectively)
```

Out[8]:

```
[array([505, 504], dtype=int32),
 array([525, 520], dtype=int32),
 array([583, 581], dtype=int32),
 array([686, 681], dtype=int32)]
```


In [9]:

```
total_no_support_vectors=[]  
for d in ds:  
    total_no_support_vectors.append(np.sum(no_support_vectors[d-1]))  
  
plt.plot(ds,total_no_support_vectors)  
plt.scatter(ds,total_no_support_vectors)  
plt.xlabel('degree')  
plt.ylabel('no. of support vectors')  
plt.show()
```



Points that lie on the margin hyperplanes should give the decision function $f(x)=0$

In [40]:


```
C=10
d=3

model=SVC(C=C,kernel='poly',degree=d)
model.fit(X_train_imputed,y_train)

#ref: https://stackoverflow.com/questions/32074239/sklearn-getting-distance-of-each
y = model.decision_function(model.support_vectors_)
pd.DataFrame(y).hist()
```

Out[40]:

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f51203a955
0>]],
      dtype=object)
```



In [42]:

```
np.sum(y==0) #there are many points that are close to the margin (f(x)=0), but there
```

Out[42]:

0

For polynomial kernel: As the parameter d (degree for polynomial kernel) increases, the margin size can be larger (more flexible) and allows more support vectors

For RBF kernel: As the parameter γ increases, the decision boundary becomes more wiggled (margin size will decrease). In addition, the no. of support vectors will increase

Q2

(a)

Ridge minimizes: $PSS + \lambda \sum \beta^2$

$$\text{error: } \sum_{i=1}^2 (y_i - \cancel{\beta_0} - \beta_1 x_{i1} - \beta_2 x_{i2})^2 + \lambda (\beta_1^2 + \beta_2^2)$$

$$= (y_1 - \beta_1 x_{11} - \beta_2 x_{12})^2 + (y_2 - \beta_1 x_{21} - \beta_2 x_{22})^2 + \lambda (\beta_1^2 + \beta_2^2)$$

$$= (y_1 - x_{11}(\beta_1 + \beta_2))^2 + (y_2 - x_{21}(\beta_1 + \beta_2))^2 + \lambda (\beta_1^2 + \beta_2^2)$$

$$= y_1^2 + y_2^2 - 2 \times (y_1 x_{11}(\beta_1 + \beta_2) + y_2 x_{21}(\beta_1 + \beta_2)) + \lambda (\beta_1^2 + \beta_2^2)$$

$$= y_1^2 + y_2^2 - 2 \times (\beta_1 + \beta_2) (y_1 x_{11} + y_2 x_{21}) + \lambda (\beta_1^2 + \beta_2^2)$$

$$= y_1^2 + y_2^2 - 2 \times \beta_1 (y_1 x_{11} + y_2 x_{21}) - 2 \times \beta_2 (y_1 x_{11} + y_2 x_{21}) + \lambda (\beta_1^2 + \beta_2^2)$$

(b)

$$\frac{\partial \text{error}}{\partial \beta_1} = 2\lambda\beta_1 - 2 \times (y_1 x_{11} + y_2 x_{21}) = 0 \Rightarrow \beta_1 = \frac{y_1 x_{11} + y_2 x_{21}}{\lambda}$$

$$\frac{\partial \text{error}}{\partial \beta_2} = 2\lambda\beta_2 - 2 \times (y_1 x_{11} + y_2 x_{21}) = 0 \Rightarrow \beta_2 = \frac{y_1 x_{11} + y_2 x_{21}}{\lambda} = \beta_1$$

(c)

Lasso minimizes $\sum \text{SS} + \lambda \sum |\beta|$

$$\text{error} = y_1^2 + y_2^2 - 2\beta_1(y_1x_{11} + y_2x_{21}) - 2\beta_2(y_1x_{12} + y_2x_{22}) + \lambda(|\beta_1| + |\beta_2|)$$

(d)

$$\begin{aligned} \frac{\partial \text{error}}{\partial \beta_1} &= \lambda \cdot \frac{|\beta_1|}{\beta_1} - 2(y_1x_{11} + y_2x_{21}) = 0 \\ \frac{\partial \text{error}}{\partial \beta_2} &= \lambda \cdot \frac{|\beta_2|}{\beta_2} - 2(y_1x_{12} + y_2x_{22}) = 0 \end{aligned} \quad \left. \begin{array}{l} \beta_1 \text{ and } \beta_2 \text{ both} \\ \text{have multiple solutions} \end{array} \right\}$$

Lasso constraint would have a diamond shape centered at the origin of the plane

Q3

In [0]:

```

from sklearn.model_selection import train_test_split

df=pd.read_csv('https://raw.githubusercontent.com/peterwu19881230/CSCE633_ML_data/main/data.csv')

print('df shape= ')
print(df.shape)

X=df.iloc[:, [0,1,2,3,4,5,6,7,8]]
response=df.iloc[:, [9]]

y=np.where(np.array(response).astype(int) == 4, 1, 0) #1 means malignant

print('no. of malignant= ' + str(np.sum(y)))
print('no. of benign= ' + str(np.sum(y==0)))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)

df shape=
(682, 10)
no. of malignant= 239
no. of benign= 443

```

There are only ~35% of the data that are malignant (imbalanced data)

Implement Decision tree

In [56]:

```
from sklearn.tree import DecisionTreeClassifier

max_depth = list(range(1,31))
criterion=['entropy','gini']
param_grid={'criterion':criterion,'max_depth':max_depth}

grid_search = GridSearchCV(DecisionTreeClassifier(), param_grid, cv=10,scoring='accu
grid_search.fit(X_train, y_train)
```

Out[56]:

```
GridSearchCV(cv=10, error_score='raise-deprecating',
             estimator=DecisionTreeClassifier(class_weight=None,
                                              criterion='gini', max_de
pth=None,
                                              max_features=None,
                                              max_leaf_nodes=None,
min_impurity_decrease=0.0,
                                              min_impurity_split=None,
                                              min_samples_leaf=1,
                                              min_samples_split=2,
min_weight_fraction_leaf=0.0,
                                              presort=False, random_st
ate=None,
                                              splitter='best'),
             iid='warn', n_jobs=None,
             param_grid={'criterion': ['entropy', 'gini'],
                          'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10,
11, 12,
                          13, 14, 15, 16, 17, 18, 19, 20,
21, 22,
                          23, 24, 25, 26, 27, 28, 29, 30]
             },
             pre_dispatch='2*n_jobs', refit=True, return_train_score=F
alse,
             scoring='accuracy', verbose=0)
```

In [60]:

```
accuracy_list=[]
for i in range(10):
    accuracy_list.append(pd.DataFrame(grid_search.cv_results_[ 'split'+str(i)+'_test_sc

accuracy_df=pd.concat(accuracy_list,axis=0)
```

```

## assign column names
params=grid_search.cv_results_['params']
params_string=['criterion:'+str(param['criterion'])+'_'+ 'max_depth:'+str(param['max_
accuracy_df.columns=params_string

mean_array=np.array([np.mean(accuracy_df[column]) for column in accuracy_df.columns
std_array=np.array([np.std(accuracy_df[column]) for column in accuracy_df.columns])

# plot based on training data
for i in [0,1]:
    degree_index=list(np.array(range(1,31))+15*i)
    plt.plot(max_depth,mean_array[degree_index])
    plt.errorbar(max_depth, mean_array[degree_index], yerr=std_array[degree_index],lo
    plt.scatter(max_depth,mean_array[degree_index])
    plt.title('criterion = '+criterion[i]+' , '+'on train')
    plt.xlabel('Max depth of the trees')
    plt.ylabel('Mean Accuracy +- std')
    plt.show()

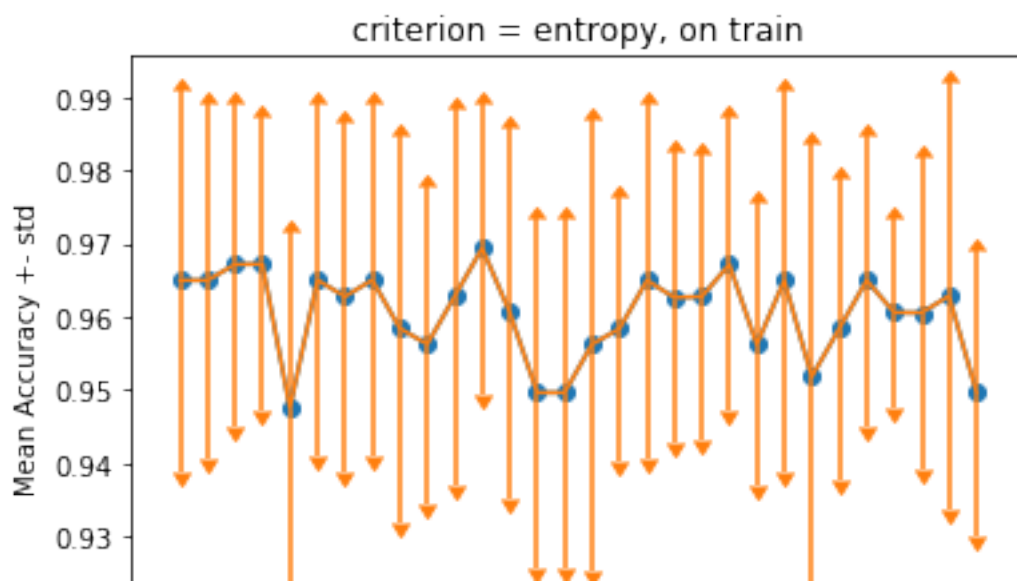
# plot based on test data


for i in [0,1]:

    accuracy_list=[]
    for j in max_depth:
        model=DecisionTreeClassifier(criterion=criterion[i],max_depth=j)
        model.fit(X_train,y_train)
        accuracy_list.append(accuracy_score(y_test,model.predict(X_test)))


    degree_index=list(np.array(range(1,31))+15*i)
    plt.plot(max_depth,np.array(accuracy_list))
    plt.plot(max_depth,np.array(accuracy_list))
    plt.title('criterion = '+criterion[i]+' , '+'on test')
    plt.xlabel('Max depth of the trees')
    plt.ylabel('Accuracy')
    plt.show()

```

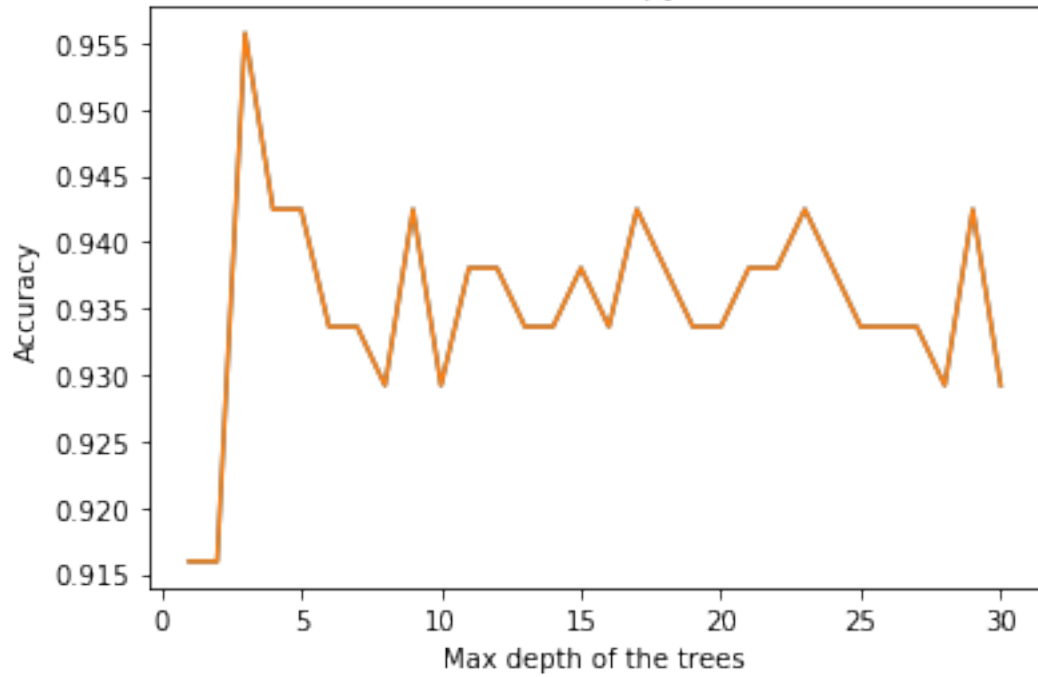





criterion = gini, on train



criterion = entropy, on test



criterion = gini, on test



max depth of the trees

Implement Random forest

In [70]:

```
from sklearn.ensemble import RandomForestClassifier

max_depth = list(range(1,31))
criterion=['entropy','gini']
param_grid={'criterion':criterion,'max_depth':max_depth}

grid_search = GridSearchCV(RandomForestClassifier(), param_grid, cv=10,scoring='accuracy',
    #Note: default no. of trees for RandomForestClassifier()= 10

grid_search.fit(X_train, np.array(y_train.ravel())) #I need .ravel() in here for RandomForestClassifier
```

Out[70]:

```
GridSearchCV(cv=10, error_score='raise-deprecating',
             estimator=RandomForestClassifier(bootstrap=True, class_weight=None,
                                               criterion='gini', max_depth=None,
                                               max_features='auto',
                                               max_leaf_nodes=None,
                                               min_impurity_decrease=0.0,
                                               min_impurity_split=None,
                                               min_samples_leaf=1,
                                               min_samples_split=2,
                                               min_weight_fraction_leaf=0.0,
                                               n_estimators=100, n_jobs=1,
                                               oob_score=False,
                                               random_state=None, verbose=0,
                                               warm_start=False),
             iid='warn', n_jobs=None,
             param_grid={'criterion': ['entropy', 'gini'],
                          'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10,
11, 12,
13, 14, 15, 16, 17, 18, 19, 20,
21, 22,
23, 24, 25, 26, 27, 28, 29, 30]
             },
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring='accuracy', verbose=0)
```

In [97]:

```
accuracy_list=[]
for i in range(10):
```

```

accuracy_list.append(pd.DataFrame(grid_search.cv_results_['split'+str(i)+'_test_score']))

accuracy_df=pd.concat(accuracy_list,axis=0)

## assign column names
params=grid_search.cv_results_['params']
params_string=['criterion:'+str(param['criterion'])+'_'+str(param['max_depth'])]
accuracy_df.columns=params_string

mean_array=np.array([np.mean(accuracy_df[column]) for column in accuracy_df.columns])
std_array=np.array([np.std(accuracy_df[column]) for column in accuracy_df.columns])

# plot based on training data
for i in [0,1]:
    degree_index=list(np.array(range(1,31))+15*i)
    plt.plot(max_depth,mean_array[degree_index])
    plt.errorbar(max_depth, mean_array[degree_index], yerr=std_array[degree_index],lo=
    plt.scatter(max_depth,mean_array[degree_index])
    plt.title('criterion = '+criterion[i]+' , '+'on train')
    plt.xlabel('Max depth of the trees')
    plt.ylabel('Mean Accuracy +- std')
    plt.show()

# plot based on test data


featureImportances=[]
for i in [0,1]:

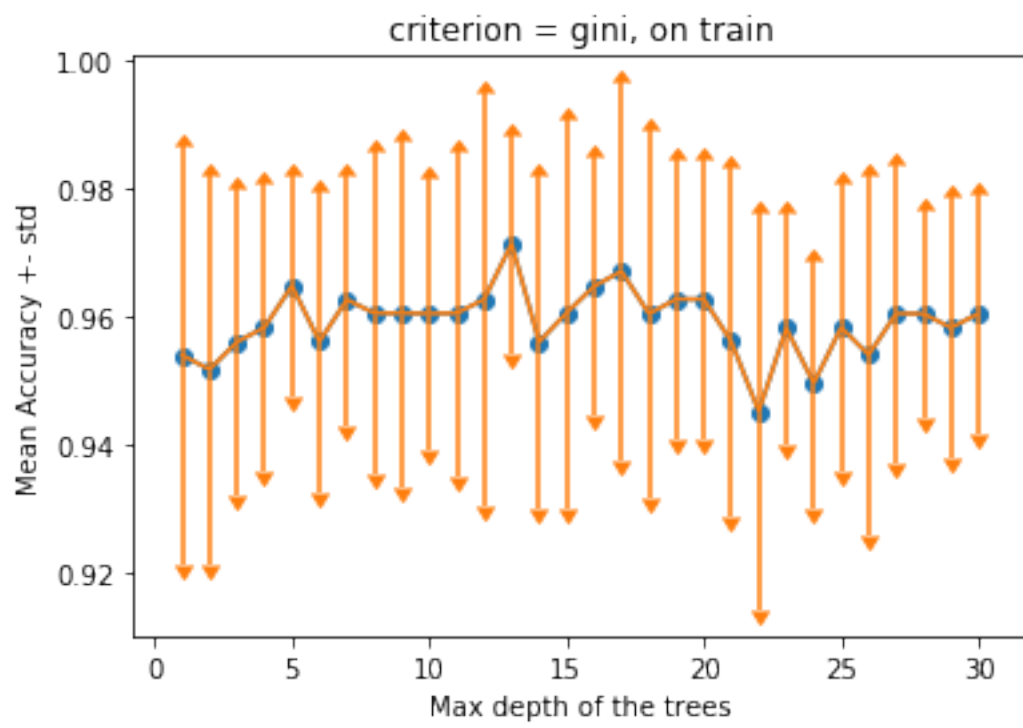

    accuracy_list=[]
    for j in max_depth:
        model=DecisionTreeClassifier(criterion=criterion[i],max_depth=j)
        model.fit(X_train,y_train)
        featureImportances.append(model.feature_importances_)
        accuracy_list.append(accuracy_score(y_test,model.predict(X_test)))


    degree_index=list(np.array(range(1,31))+15*i)
    plt.plot(max_depth,np.array(accuracy_list))
    plt.plot(max_depth,np.array(accuracy_list))
    plt.title('criterion = '+criterion[i]+' , '+'on test')
    plt.xlabel('Max depth of the trees')
    plt.ylabel('Accuracy')
    plt.show()

```

criterion = entropy, on train







```
In [143]:  
  
# show feature importance rankings (using average of all 60 models based on [gini,ex  
features=['Clump Thickness','Uniformity of Cell Size','Uniformity of Cell Shape','Ma  
  
import scipy.stats as ss  
def rank(feature_importance):  
    ranking=len(feature_importance) - ss.rankdata(feature_importance,method="max").ast  
    return(ranking)  
  
ranking_list=[]  
importance_list=[]  
for i in range(60):  
    importance_list.append(pd.DataFrame(featureImportances[i]).T)  
    ranking_list.append(pd.DataFrame(rank(featureImportances[i])).T)  
  
ranking_df=pd.concat(ranking_list,axis=0)  
importance_df=pd.concat(importance_list,axis=0)  
feature_ranking_list=[pd.DataFrame(features),pd.DataFrame(importance_df.mean(axis=0)  
  
feature_ranking_df=pd.DataFrame(pd.concat(feature_ranking_list,axis=1))  
feature_ranking_df.columns=['features','importance mean','ranking mean','ranking std  
print(feature_ranking_df)
```

	features	importance mean	ranking mean	ranking
std				
0	Clump Thickness	0.111992	2.033333	0.18
1020				
1	Uniformity of Cell Size	0.767551	1.000000	0.00
0000				
2	Uniformity of Cell Shape	0.015073	4.766667	1.31
9562				
3	Marginal Adhesion	0.004198	7.083333	1.57
6218				
4	Single Epithelial Cell Size	0.010265	5.616667	1.15
1148				
5	Bare Nuclei	0.068001	2.933333	0.25
1549				
6	Bland Chromatin	0.012347	5.216667	1.48
5428				
7	Normal Nucleoli	0.008863	6.116667	1.82
3497				
8	Mitoses	0.001711	7.716667	1.76
6848				

The rankings were based on the importance of each feature averaged by all the 60 models (by [gini,entropy] X [1~30 max depth of tree]). The accuracy doesn't change significantly across different hyperparameters so averaging all of them should be fine. For the final model I would probably take out 'Marginal Adhesion' and 'Mitosis' because they are the least important. This can be further confirmed by calculating accuracy after post-pruning the trees

Homework 2

This jupyter notebook is also accessble as an online colab notebook :

<https://colab.research.google.com/drive/10KcuDSp9jCx0eFkdQ-UvrS6VPQjmpYPQ#scrollTo=Yj8OAlyA5Tzv>
(<https://colab.research.google.com/drive/10KcuDSp9jCx0eFkdQ-UvrS6VPQjmpYPQ#scrollTo=Yj8OAlyA5Tzv>)

Q1

Load & normalize the data