

# CSCE 633: Machine Learning

## Lecture 8: Model Selection and Regularization

Texas A&M University

9-11-18

# Last Time

- Logistic Regression
- Gradient Descent

# Goals of this lecture

- Ridge Regression
- Lasso Regularization
- Measures of Model Performance

## Shrinkage

- Train a model fitting all  $p$  predictors that constrains or regularizes the coefficients
- Two best methods for this, Ridge Regression and the Lasso

## Ridge Regression

- Recall the least squares fit for  $\beta_0, \beta_1, \dots, \beta_p$  minimizes
- $RSS = \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2$

## Ridge Regression

- Recall the least squares fit for  $\beta_0, \beta_1, \dots, \beta_p$  minimizes
- $RSS = \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2$
- With Ridge Regression, we modify the equation that we want to minimize by adding a penalty (paying a price) for using predictors

## Ridge Regression

- Recall the least squares fit for  $\beta_0, \beta_1, \dots, \beta_p$  minimizes
- $RSS = \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2$
- With Ridge Regression, we modify the equation that we want to minimize by adding a penalty (paying a price) for using predictors

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 = RSS + \lambda \sum_{j=1}^p \beta_j^2$$

- Where  $\lambda \geq 0$  is the tuning parameter

## Ridge Regression

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 = RSS + \lambda \sum_{j=1}^p \beta_j^2$$

- Where  $\lambda \geq 0$  is the tuning parameter
- Ridge Regression creates a tradeoff, we still want coefficients that reduce  $RSS$  but now we have a shrinkage penalty.
- This penalty is small if  $\beta_0, \dots, \beta_p$  are close to 0
- Where least squares creates a single set of coefficients, Ridge Regression creates a set  $\hat{\beta}_\lambda^R$  for each  $\lambda$




## Ridge Regression

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 = RSS + \lambda \sum_{j=1}^p \beta_j^2$$


- Where least squares creates a single set of coefficients, Ridge Regression creates a set  $\hat{\beta}_\lambda^R$  for each  $\lambda$
- Selecting the right  $\lambda$  is key (done via cross-validation)
- Note, penalty is not assigned to the intercept  $\beta_0$  since  $\beta_0$  is the measure of the mean value of the response when  $x_{i1} = x_{i2} = \dots = x_{ip} = 0$
- If we assume columns of  $X$  have been centered (meaning each column has a mean of 0) then intercept is  $\hat{\beta}_0 = \bar{y} = \sum_{i=1}^n \frac{y_i}{n}$

## Ridge Regression and Credit Data




- Each line is one of the ten variables as a function of  $\lambda$ . Solid black line is the income variable
- We can see when  $\lambda = 0$  we get standard least squares.
- When  $\lambda = \infty$  we have the null model

## Ridge Regression and Credit Data



- Income, limit, rating, and student have the largest coefficients.
- Note, in some steps individual estimates might actually grow because of relative importance
- Right hand side of figure, we scale the x-axis as  $\frac{\|\hat{\beta}_\lambda^R\|_2}{\|\hat{\beta}\|_2}$  where  $\|\hat{\beta}\|_2$  is the  $\ell_2$  norm of the least squares coefficient estimates

## Ridge Regression and Credit Data



- Right hand side of figure, we scale the x-axis as  $\frac{\|\hat{\beta}_\lambda^R\|_2}{\|\hat{\beta}\|_2}$  where  $\|\hat{\beta}\|_2$  is the  $\ell_2$  norm of the least squares coefficient estimates
- That value ranges from 1 when  $\lambda = 0$  to 0 when  $\lambda = \infty$
- Thus, the x-axis represents amounts coefficient estimates have been shrunk to 0

## Scaling

- Scaling is now an important part of data we need to consider.
- Whereas, with least squares, if  $X_j$  was scaled by some constant  $c$  then least squares  $\beta$  would have been scaled by  $\frac{1}{c}$ , this is not true for Ridge Regression
- $x_j \hat{\beta}_{j,\lambda}^R$  will depend on  $\lambda$  and scaling of  $x_j$  and perhaps even the scaling of other predictors.
- To avoid scaling issues, we should standardize predictors

$$\tilde{x}_{ij} = \frac{x_{ij}}{\sqrt{\frac{1}{n} \sum_{i=1}^n (x_{ij} - \bar{x}_{ij})^2}}$$

- Where the denominator is the estimated standard deviation of  $j$


## Normalization

- One important step in numeric data is normalization of the data to help these techniques
- one common technique is to center and scale each predictor  $X_j$
- This causes all the predictors to have mean 0 and standard deviation of 1

$$\tilde{x}_j = \frac{x_j - \bar{x}_j}{\sigma_j}$$

## Why does this work?

- Rooted in the bias-variance trade-off
- As  $\lambda$  increases, flexibility of ridge regression fit decreases, decreasing variance but increasing bias.



- Simulated data of  $p = 45$ ,  $n = 50$ , black is bias, green is variance, purple is test error
- $\lambda = 30$  is the optimal solution and  $MSE$  of least squares is almost as high as null-model

## LASSO

- Ridge Regression has one obvious disadvantage. Unlike subset methods, ridge regression still fits all  $p$  predictors.
- The penalty  $\lambda \sum_j \beta_j^2$  will shrink all coefficients but none will hit 0 exactly
- This may not be a problem for accuracy but is for interpretability
- For example with our credit data set, Ridge Regression will still use all 10 predictors, even if it finds that income, limit, rating, and student are the most important.

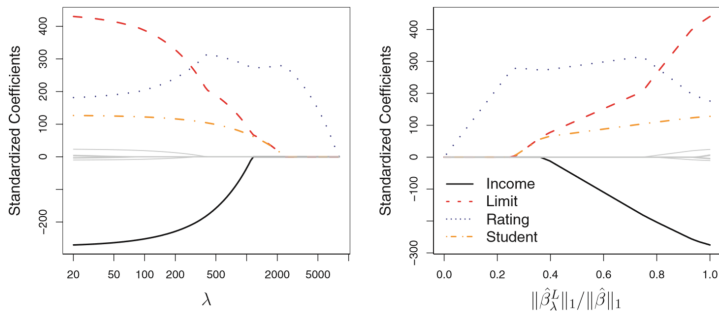


## LASSO Regularization

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p |\beta_j| = RSS + \lambda \sum_{j=1}^p |\beta_j|$$

- Creates a set  $\hat{\beta}_{\lambda}^L$  for each  $\lambda$
- We use the  $\ell_1$  norm instead of  $\ell_2$
- Lasso shrinks coefficients but the  $\ell_1$  penalty drives coefficients to 0 when  $\lambda$  is sufficiently large
- This means Lasso performs variable selection!

## Lasso and Credit Data



- Lasso picks rating, then student and limit together, then income. Eventually all others would enter as you approach least squares fit
- Where ridge selects coefficients/shrinkage, lasso produces models with any number of variables

## Another Formulation


$$\min_{\beta} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2$$

Subject to  $\sum_{j=1}^p |\beta_j| \leq s$  for Lasso

and Subject to  $\sum_{j=1}^p \beta_j^2 \leq s$  for Ridge Regression

- If  $p = 2$  Lasso solution falls within the diamond  $|\beta_1| + |\beta_2| \leq s$
- Circle for Ridge  $\beta_1^2 + \beta_2^2 \leq s$



## Another Formulation




- Ellipses are increasing  $RSS$  from the least squares solution
- if the  $\lambda$  allows enough to include  $RSS$  that is the fit found
- Because Lasso will intersect at a corner, while Ridge somewhere in between on circle, is why Lasso sets some coefficients to 0 while Ridge just shrinks them

## Final Notes


- Lasso is better if small set of predictors dominates response
- Ridge is better if all predictors contribute somewhat equally
- Cannot tell in advance, need cross-validation to give us an idea
- Lasso shrinks very differently than Ridge, known as soft thresholding
- Ridge assumes the density function of the posterior probabilities of  $\beta$  are Gaussian (most coefficients are somewhere near 0), while Lasso assumes Laplacian (most coefficients centered at 0)



## More about sparsity



$$0.5 \times (x-v)^2 + \lambda |x|$$



$$0.5 \times (x-v)^2 + \lambda x^2$$

If  $v \geq \lambda$ ,  $x = v - \lambda$

If  $v \leq -\lambda$ ,  $x = v + \lambda$



Else,  $x = 0$

$$x = v / (1 + 2\lambda)$$

Nondifferentiable at 0

Differentiable at 0

## More about sparsity



The gradient of the  $\ell_2$ -norm vanishes when  $\alpha$  get close to 0. On its differentiable part, the norm of the gradient of the  $\ell_1$ -norm is constant.

## How to Solve LASSO

Rewrite the optimization problem:


$$\min_{\beta} \frac{1}{2} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1$$

Challenges:

- The optimization is non-smooth.
- Subgradient Method
  - Subgradients are easy to derive and implement
  - Convergence needs carefully chosen step sizes
  - Convergence is weak theoretically



## Subgradient Method



## How to Solve LASSO

Fast  $l_1$  minimization algorithms:

- Iterative Shrinkage Thresholding Algorithm (ISTA)
- Proximal Gradient Method (PGM)
- Alternating Direction Methods of Multipliers

## Iterative Shrinkage Thresholding Algorithm (ISTA)

ISTA considers the LASSO model as a special case of the composite objective function:

$$\min_{\beta} F(\beta) = f(\beta) + g(\beta),$$

where  $f$  is a smooth and convex function, and  $g$  is the regularization term that is not necessarily smooth nor convex. Here  $f(\beta) = \frac{1}{2} \|y - X\beta\|_2^2$ .

- If  $g(\beta) = \lambda \|\beta\|_2^2$ : Ridge regression.
- If  $g(\beta) = \lambda \|\beta\|_1$ : LASSO.

## ISTA using Hessian Matrix Approximation

Estimate  $f(\beta)$  using its Taylor expansion to the second order around  $\beta^k$ :

$$\begin{aligned}\beta^{k+1} &= \operatorname{argmin}_{\beta} \left\{ f(\beta^k) + \nabla f(\beta^k)^T (\beta - \beta^k) + \frac{1}{2} (\beta - \beta^k)^T \nabla^2 f(\beta^k) (\beta - \beta^k) + g(\beta) \right\} \\ &\approx \operatorname{argmin}_{\beta} \left\{ \nabla f(\beta^k)^T (\beta - \beta^k) + \frac{\alpha^k}{2} \|\beta - \beta^k\|_2^2 + g(\beta) \right\} \\ &= \operatorname{argmin}_{\beta} \left\{ \frac{\alpha^k}{2} \|\beta - \gamma^k\|_2^2 + g(\beta) \right\}\end{aligned}$$

where

$$\gamma^k = \beta^k - \frac{1}{\alpha^k} \nabla f(\beta^k).$$

## ISTA using Hessian Matrix Approximation

Specifically for LASSO, where  $g(\beta) = \|\beta\|_1$ , the last optimization step is separable:


$$\beta^{k+1} = \operatorname{argmin}_{\beta} \left\{ \sum_i \frac{\alpha^k}{2} (\beta_i - \gamma_i^k)^2 + \lambda |\beta_i| \right\}$$

The problem consists of multiple independent 1-D problems that have explicit solution:


$$\beta_i^{k+1} = \operatorname{soft}(\gamma_i^k, \frac{\lambda}{\alpha^k})$$

$$\begin{aligned} \operatorname{soft}(u, a) &\doteq \operatorname{sgn}(u) \max\{|u| - a, 0\} \\ &= \begin{cases} \operatorname{sgn}(u)(|u| - a) & \text{if } |u| > a \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

## Soft Thresholding Function



(a) Soft-thresholding operator,  
 $\alpha^{\text{st}} = \text{sign}(\beta) \max(|\beta| - \lambda, 0)$ .



(b) Hard-thresholding operator  
 $\alpha^{\text{ht}} = 1_{|\beta| \geq \mu} \beta$ .

## ISTA using Proximal Gradient Method

At each step, perform a gradient descent step on  $f(\beta)$  without considering the non-smooth regularization:

$$\gamma^k = \beta^k - \alpha^k \nabla f(\beta^k) = \beta^k + \alpha^k X^T (y - X\beta^k)$$

Then combine the regularization by solving the following proximity problem:

$$\beta^{k+1} = \underset{\beta}{\operatorname{argmin}} \left\{ \frac{1}{2\alpha^k} \|\beta - \gamma^k\|_2^2 + \lambda \|\beta\|_1 \right\}$$

which will induce exactly the same solution:

$$\beta^{k+1} = \operatorname{soft}(\beta^k + \alpha^k X^T (y - X\beta^k), \alpha^k \lambda).$$

Need select a small enough step size  $\alpha^k$ .

## Lasso: Continued


- Selecting  $\lambda$
- Further Lasso Solvers and Elastic Net




## LASSO: Picking $\lambda$

- Need to pick best  $\lambda$  (or  $s$  in the alternative formulation) for best estimation
- We can run a cross-validation over a grid of  $\lambda$  values
- We pick the *lambda* with the smallest error

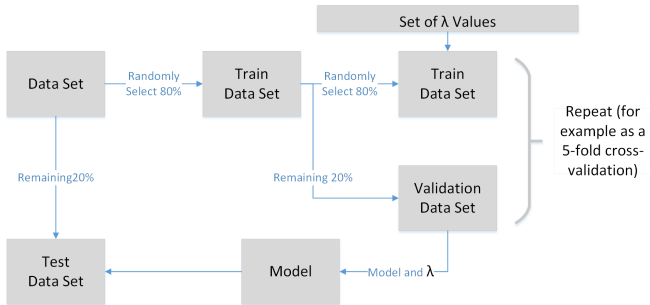
## LASSO: Picking $\lambda$




## LASSO: Picking $\lambda$



# LASSO: Picking $\lambda$




## LASSO: Our Credit Example



- Sometimes Lasso does not do better than Least Squares Solution
- Small  $\lambda$  selected here

## LASSO: Synthetic Example



- Sometimes Lasso does a lot better than Least Squares Solution

## LASSO: Elastic Net

- Last week, we discussed the differences between Ridge Regression and Lasso
- We also discussed how it is not immediately obvious which would be better, sometimes need cross-validation to test
- if  $n > p$  but variables are correlated, ridge empirically does better than lasso
- if  $p > n$  lasso cannot select more than  $n$  variables before it saturates.
- A mix between Lasso and Ridge exists, called Elastic Net

## Vanilla Elastic Net

New Objective Function is

$$J(\beta, \lambda_1, \lambda_2) = \|y - X\beta\|^2 + \lambda_2 \|\beta\|_2^2 + \lambda_1 \|\beta\|_1$$

- The objective now has a penalty that is from ridge regression and a penalty that is from lasso
- It turns out this doesn't predict really well, unless the optimal solution is found by ridge or by lasso
- This is because some solution in the middle has coefficients penalized by both  $\lambda_1$  and  $\lambda_2$



## Vanilla Elastic Net

New Objective Function is

$$J(\beta, \lambda_1, \lambda_2) = \|y - X\beta\|^2 + \lambda_2\|\beta\|_2^2 + \lambda_1\|\beta\|_1$$

- The objective now has a penalty that is from ridge regression and a penalty that is from lasso
- It turns out this doesn't predict really well, unless the optimal solution is found by ridge or by lasso
- This is because some solution in the middle has coefficients penalized by both  $\lambda_1$  and  $\lambda_2$
- To fix it, we adjust the optimal solution. So, first, we solve the vanilla version

## LARS-Elastic Net

First we re-write  $X$  as

$$\tilde{X} = \frac{1}{\sqrt{1 + \lambda_2}} \begin{pmatrix} X \\ \sqrt{\lambda_2} I_p \end{pmatrix}$$

Where  $I_p$  is the identity matrix and

$$\tilde{y} = \begin{pmatrix} y \\ 0_{p \times 1} \end{pmatrix}$$

Then we solve for  $\beta$  like a normal lasso problem

$$\tilde{\beta} = \underset{\tilde{\beta}}{\operatorname{argmin}} \|\tilde{y} - \tilde{X}\tilde{\beta}\|^2 + \frac{\lambda_1}{\sqrt{1 + \lambda_2}} \|\tilde{\beta}\|_1$$

$$\text{So } \beta = \frac{\tilde{\beta}}{\sqrt{1 + \lambda_2}}$$

## Improved Elastic Net

Then we solve for  $\beta$  like a normal lasso problem

$$\tilde{\beta} = \operatorname{argmin}_{\tilde{\beta}} \|\tilde{y} - \tilde{X}\tilde{\beta}\|^2 + \frac{\lambda_1}{\sqrt{1 + \lambda_2}} \|\tilde{\beta}\|_1$$

So  $\beta = \frac{\tilde{\beta}}{\sqrt{1 + \lambda_2}}$

- So now we want to undo one of the penalties so coefficients aren't double penalized
- for simplicity we undo the  $\lambda_2$  penalty ( $\ell_2$ )

$$\hat{\beta} = \sqrt{1 + \lambda_2} \tilde{\beta}$$

# Measures of Classification Accuracy

- Confusion Matrices
- Class labels versus Probabilities
- Measurements

## Confusion Matrix

		<i>True default status</i>		
		No	Yes	Total
<i>Predicted default status</i>	No	9,644	252	9,896
	Yes	23	81	104
Total		9,667	333	10,000

- If we denote true positives as  $TP$ , false positives as  $FP$ , true negatives as  $TN$ , and false negatives as  $FN$  then we can say Accuracy is
- $ACC = \frac{TP+TN}{TP+FN+TN+FP}$
- What is the accuracy here?

## Confusion Matrix

		<i>True default status</i>		
		No	Yes	Total
<i>Predicted default status</i>	No	9,644	252	9,896
	Yes	23	81	104
Total		9,667	333	10,000

- If we denote true positives as  $TP$ , false positives as  $FP$ , true negatives as  $TN$ , and false negatives as  $FN$  then we can say Accuracy is
- $ACC = \frac{TP+TN}{TP+FN+TN+FP}$
- What is the accuracy here? 97.25%

## Confusion Matrix: Problems with Accuracy?

		True	
		No	Yes
Predicted	No	50	50
	Yes	0	0

- If we denote true positives as  $TP$ , false positives as  $FP$ , true negatives as  $TN$ , and false negatives as  $FN$  then we can say Accuracy is
- $ACC = \frac{TP+TN}{TP+FN+TN+FP}$

## Confusion Matrix with Class Imbalance: Problems with Accuracy?

		True	
		No	Yes
Predicted	No	900	100
	Yes	0	0

- If we denote true positives as  $TP$ , false positives as  $FP$ , true negatives as  $TN$ , and false negatives as  $FN$  then we can say Accuracy is
- $ACC = \frac{TP+TN}{TP+FN+TN+FP}$




## Confusion Matrix: Measurements in Addition to Accuracy

		True default status		
		No	Yes	Total
Predicted default status	No	9,644	252	9,896
	Yes	23	81	104
Total		9,667	333	10,000

- If we denote true positives as  $TP$ , false positives as  $FP$  (Type I Error), true negatives as  $TN$ , and false negatives as  $FN$  (Type II Error) then we can say Accuracy is
- $ACC = \frac{TP+TN}{TP+FN+TN+FP}$
- Recall = Sensitivity = True Positive Rate =  $\frac{TP}{TP+FN}$
- Specificity = True Negative Rate =  $\frac{TN}{TN+FP} = 1 - FPR$
- False Positive Rate (FPR) =  $\frac{FP}{FP+TN}$
- Precision = Positive Predictive Value =  $\frac{TP}{TP+FP}$
- F1 Score =  $2 \times \frac{Precision \times Recall}{Precision + Recall}$
- Can calculate these per-class and average together or total across all classes

## Logistic Regression: Decision Threshold



## Logistic Regression: Decision Threshold

If we pick a decision threshold of  $p(x) > 0.5$  what happens?

Predicted	Ground Truth
2%	0
3%	0
5%	1
1%	0
15%	0
25%	1
24%	1
13%	0
8%	0
12%	1

## Logistic Regression: Decision Threshold


Predicted	TRUE	50	25	24	15	13	12	8	5	3	2	1
1	0	0	0	0	0	0	0	0	0	0	0	1
2	0	0	0	0	0	0	0	0	0	0	1	1
3	0	0	0	0	0	0	0	0	0	1	1	1
5	1	0	0	0	0	0	0	0	1	1	1	1
8	0	0	0	0	0	0	0	1	1	1	1	1
12	1	0	0	0	0	0	1	1	1	1	1	1
13	0	0	0	0	0	1	1	1	1	1	1	1
15	0	0	0	0	1	1	1	1	1	1	1	1
24	1	0	0	1	1	1	1	1	1	1	1	1
25	1	0	1	1	1	1	1	1	1	1	1	1

## Logistic Regression: Decision Threshold


	50	25	24	15	13	12	8	5	3	2	1
TP	0	1	2	2	2	3	3	4	4	4	4
FP	0	0	0	1	2	2	3	3	4	5	6
TN	6	6	6	5	4	4	3	3	2	1	0
FN	4	3	2	2	2	1	1	0	0	0	0
TPR	0	0.25	0.5	0.5	0.5	0.75	0.75	1	1	1	1
FPR	0	0	0	0.166667	0.333333	0.333333	0.5	0.5	0.666667	0.833333	1

## Logistic Regression: Decision Threshold


	50	25	24	15	13	12	8	5	3	2	1
TP	0	1	2	2	2	3	3	4	4	4	4
FP	0	0	0	1	2	2	3	3	4	5	6
TN	6	6	6	5	4	4	3	3	2	1	0
FN	4	3	2	2	2	1	1	0	0	0	0
TPR	0	0.25	0.5	0.5	0.5	0.75	0.75	1	1	1	1
FPR	0	0	0	0.166667	0.333333	0.333333	0.5	0.5	0.666667	0.833333	1



# Receiver Operating Characteristic Curve



## ROC and Measurements



- By predicting probability instead of label can generate ROC Curve
- Can choose threshold on ROC curve that optimizes some threshold-specific Measurement
- Can also plot Precision-Recall Curve
- Area Under ROC Curve (AUROC) is a measurement of how well your model separates classes (without taking decision threshold into account)
- Can also calculate AUPRC



## Problem Solving: Gradient Descent

$$MSE(\beta_0, \beta_1) = J(\beta_0, \beta_1) = J(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_i))^2$$

- Optimize  $\beta_0$  and  $\beta_1$  by Gradient Descent
- Remember, Update Rules  $\beta_0^{t+1} = \beta_0^t - \alpha \frac{\partial}{\partial \beta_0} J(\beta_0, \beta_1)$
- Remember, Update Rules  $\beta_1^{t+1} = \beta_1^t - \alpha \frac{\partial}{\partial \beta_1} J(\beta_0, \beta_1)$
- if  $\beta = (\beta_0, \beta_1)$  then
- $\beta^{t+1} = \beta^k - \alpha \nabla J(\beta)$

## Problem Solving: Jacobian

$$\nabla J(\beta)$$

- If we have a series of functions  $f$  that map input  $x$  to an output  $y$  (such as linear regression)
- The Jacobian is a matrix such that  $J_{ij} = \frac{\partial f_i}{\partial x_j}$
- The Hessian is a matrix such that  $H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$

## Problem Solving: Jacobian

$$f(x, y) = \begin{bmatrix} x^2y \\ 5x + \sin(y) \end{bmatrix}$$

$$\nabla f = \begin{bmatrix} \frac{\partial f_1}{\partial x} & \frac{\partial f_1}{\partial y} \\ \frac{\partial f_2}{\partial x} & \frac{\partial f_2}{\partial y} \end{bmatrix} = \begin{bmatrix} 2xy & x^2 \\ 5 & \cos(y) \end{bmatrix}$$

## Problem Solving: Gradient Descent

$$MSE(\beta_0, \beta_1) = J(\beta_0, \beta_1) = J(\beta) = \frac{1}{2n} \sum_{i=1}^n ((\beta_0 + \beta_1 x_i) - y_i)^2$$

Our Objective

$$\min_{\beta_0, \beta_1} J(\beta_0, \beta_1) = \min_{\beta} J(\beta)$$

Then our update rule at time  $t + 1$ , with stopping criteria  $\epsilon$  is

$$\beta^{t+1} = \beta^t - \alpha \nabla J(\beta)$$

$$\nabla J(\beta) = \left[ \frac{\partial J(\beta)}{\partial \beta_0} \quad \frac{\partial J(\beta)}{\partial \beta_1} \right]$$

$$\|\nabla J(\beta)\| < \epsilon$$

## Problem Solving: Gradient Descent

$$\nabla J(\beta) = \begin{bmatrix} \frac{\partial J(\beta)}{\partial \beta_0} & \frac{\partial J(\beta)}{\partial \beta_1} \end{bmatrix}$$

$$= \begin{bmatrix} \frac{1}{n} \sum_{i=1}^n (\beta_0 + \beta_1 x_i - y_i) & \frac{1}{n} \sum_{i=1}^n (\beta_0 + \beta_1 x_i - y_i) x_i \end{bmatrix}$$

# Takeaways and Next Time

- Model selection finds optimal solutions by using a subset of predictors
- Regularization minimizes RSS subject to a price for the predictors used
- Model Performance and Problem Solving
- Next Time: Measures of Performance with Logistic Regression